

In [1]:



```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn import preprocessing, svm
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
```

In [2]:



```
df=pd.read_csv(r"c:\Users\DELL\Documents\bottle.csv")  
df
```

C:\Users\DELL\AppData\Local\Temp\ipykernel_5452\3296139706.py:1: DtypeWarning: Columns (47,73) have mixed types. Specify dtype option on import or set low_memory=False.
df=pd.read_csv(r"c:\Users\DELL\Documents\bottle.csv")

Out[2]:

Cst_Cnt	Btl_Cnt	Sta_ID	Depth_ID	Depthm	T_degC	Salnty	O2ml_L	STheta	(
0	1	1	054.0 056.0	19-4903CR-HY-060-0930-05400560-0000A-3	0	10.500	33.4400	NaN	25.64900
1	1	2	054.0 056.0	19-4903CR-HY-060-0930-05400560-0008A-3	8	10.460	33.4400	NaN	25.65600

In [3]:

df=df[['Salnty','T_degC']]

In [4]:

df.columns=['Sal','Temp']

df

Out[4]:

	Sal	Temp							
0	33.4400	10.500							
4	33.4400	10.460	5	054.0	19-4903CR-HY-060-0930-05400560-0020A-7	20	10.450	33.4210	NaN 25.64300
2	33.4370	10.460							
3	33.4200	10.450...	
4	33.4210	10.450							
864858	34404	864859	093.4	026.4	20-1611SR-MX-310-2239-09340264-0000A-7	0	18.744	33.4083	5.805 23.87055 1
864858	33.4083	18.744							
864859	33.4083	18.744							
864860	33.4150	18.692							
864861	33.4062	18.692	093.4	026.4	20-1611SR-MX-310-2239-09340264-0002A-3	2	18.744	33.4083	5.805 23.87072 1
864862	33.3880	17.533							
864863 rows × 2 columns									
864860	34404	864861	093.4	026.4	20-1611SR-MX-310-2239-09340264-0005A-3	5	18.692	33.4150	5.796 23.88911 1
864861	34404	864862	093.4	026.4	20-1611SR-MX-310-2239-09340264-0010A-3	10	18.161	33.4062	5.816 24.01426 1

In [5]:

Cst_Cnt Btl_Cnt Sta_ID Depth_ID Depthm T_degC Salnty O2ml_L STheta (

df.head(10)

				20-							
				1611SR-							
Out[5]:	34404	864863	093.4	MX-310-	15	17.533	33.3880	5.774	24.15297	1	
864862			026.4	2239-							
	Sal	Temp		09340264-							
				0015A-3							
0	33.440	10.50									

864863 rows x 11 columns

2	33.437	10.46
3	33.420	10.45
4	33.421	10.45
5	33.431	10.45
6	33.440	10.45
7	33.424	10.24
8	33.420	10.06
9	33.494	9.86

In [6]:

df.describe()

Out[6]:

	Sal	Temp
count	817509.000000	853900.000000
mean	33.840350	10.799677
std	0.461843	4.243825
min	28.431000	1.440000
25%	33.488000	7.680000
50%	33.863000	10.060000
75%	34.196900	13.880000
max	37.034000	31.140000

In [7]:

df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 864863 entries, 0 to 864862
Data columns (total 2 columns):
Column Non-Null Count Dtype
--- -
0 Sal 817509 non-null float64
1 Temp 853900 non-null float64
dtypes: float64(2)
memory usage: 13.2 MB

In [8]:

```
df.fillna(method='ffill',inplace=True)
```

C:\Users\DELL\AppData\Local\Temp\ipykernel_5452\4116506308.py:1: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
df.fillna(method='ffill',inplace=True)
```

In [12]:

```
x=np.array(df['Sal']).reshape(-1,1)  
y=np.array(df['Temp']).reshape(-1,1)
```

In [13]:

```
df.dropna(inplace=True)
```

C:\Users\DELL\AppData\Local\Temp\ipykernel_5452\1379821321.py:1: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
df.dropna(inplace=True)
```


In [14]:

```
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.25)  
regr=LinearRegression()  
regr.fit(x_train,y_train)  
print(regr.score(x_test,y_test))
```

0.20593523576030115

In [16]:

```
y_pred=regr.predict(x_test)  
plt.scatter(x_test,y_test,color='b')  
plt.plot(x_test,y_pred,color='k')  
plt.show()
```




In [21]:

```
df500=df[:][:500]  
sns.lmplot(x="Sal",y="Temp",data=df500,order=1,ci=None)
```

Out[21]:


<seaborn.axisgrid.FacetGrid at 0x20214030670>



In [23]:

```
df500.fillna(method='ffill',inplace=True)
x=np.array(df500['Sal']).reshape(-1,1)
y=np.array(df500['Temp']).reshape(-1,1)
df500.dropna(inplace=True)
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.25)
regr=LinearRegression()
regr.fit(x_train,y_train)
print("Regression:",regr.score(x_test,y_test))
y_pred=regr.predict(x_test)
plt.scatter(x_test,y_test,color='b')
plt.plot(x_test,y_pred,color='k')
plt.show()
```

Regression: 0.8227647198046334



In [25]:

```
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score
model=LinearRegression()
model.fit(x_train,y_train)
y_pred=model.predict(x_test)
r2=r2_score(y_test,y_pred)
print("R2 score:",r2)
```

R2 score: 0.8227647198046334

In []:

